

Pranav Pappu

pranavpappu@u.nus.edu | +65 9754 0582 | LinkedIn | GitHub | Website | X

Education

National University of Singapore

2025 – 2029

Bachelor of Computing (Honours) in Computer Science

AIE Scholarship – Selected as 1 of 20 recipients

Technical Skills

Languages/Runtime: Zig • C++ • Python • TypeScript/JavaScript • Rust • Node.js • Bun

ML/LLM: PyTorch • GPT-2/Transformers • Hugging Face • tiktoken • OpenAI • Gemini • FunctionGemma

Frameworks/Tooling: React • React Native/Expo • FastAPI/Uvicorn • Tauri • Vite/Tailwind • Playwright

Data/Protocol/Integrity: MCP (stdio/HTTP) • JSON/JSONL/Schema • Zig Build • SHA-256/HMAC-SHA-256

Experience

Founder, Tweakler

2026 – Present

- Founded Tweakler, a reference-led website generation system that helps agents avoid generic UI: from a one-line brief, it derives an audience's values and jobs-to-be-done, surfaces diverse permissioned design references, and transforms a selected reference into a distinct, editable site with brand-specific copy, assets, and visual direction. Piloted with studios including [Brightsea Studio](#).
- Built a [54-tool MCP workflow](#) and deterministic browser capture-to-code engine for rights-gated cloning, responsive site generation, transactional copy and asset replacement with rollback, and fail-closed desktop/mobile visual, accessibility, build, and type-check validation.

Research Assistant, National University of Singapore

Jan 2026 – Present

- Built a dynamic 3DGS benchmark with [11 YAML pipelines](#) across [6 method repos](#) and [3 datasets](#), plus a [6-metric evaluator](#) for repeatable quality/performance reports.

CCSGP Fellow (AssessMate), NUS Centre for Computing for Social Good

2025 – Present

- Building AssessMate, an AI case-notes copilot for structured documentation; piloted with [2 assessors](#) across multiple sessions to measure documentation-time reduction on PDPA-sensitive child-development data.

C++ and Python Specialist (Freelance), Outlier.ai – Remote

Dec 2024 – May 2025

- Evaluated C++/Python model outputs for correctness, edge-case coverage, and reasoning quality, contributing feedback for code-generation model improvement.

Projects

EvoServe – Harness Evaluation & Quality Control

2026 – Present

- Built EvoServe, a harness-quality and evaluation control plane for coding-agent systems: compares policy candidates, ingests captured runs, independently verifies evidence, and produces auditable reports for quality and regression review.
- Implemented a [10-command](#) validate/replay/score/verify/report CLI with [36 JSON schemas](#) and [24 Zig tests](#); fails closed on malformed, missing, ambiguous, mutated, or untrusted evidence while preserving explicit review before policy activation.

GPT-2 (124M) Recreation from Scratch

2024

- Recreated GPT-2 124M from scratch (tokenizer, transformer blocks, optimizer, mixed precision, and FineWeb Edu evaluation); trained on [10B tokens](#).

Remember – Memory-Augmented AR Assistant

2026

- Built the [1st-place Replit Agent Hackathon](#) project, an AR memory assistant integrating [4 AI workflows](#) for real-time face recognition, speech transcription, LLM memory extraction, and voice Q&A; linked [128-float face embeddings](#) to a persistent personal knowledge graph for context-aware recall across repeated encounters.

Competitions

Replit Agent Hackathon

2026

- Won [1st place](#) with Remember, a browser-based smart-glasses assistant helping users recognize people, preserve conversational context, and retrieve relevant memories through voice.

Google DeepMind × Cactus (YC S25) Global Hackathon

2026

- [2nd place](#): Built SecureClaw, a privacy router for [12 PII types](#) with [3 routing modes](#), on-device/cloud inference, and a [30-case benchmark](#).