

Pranav Pappu

pranavpappu@u.nus.edu | +65 9754 0582 | LinkedIn | GitHub | pranavpappu.dev | X

Availability: Jan–Dec 2027 for internships in the San Francisco Bay Area

Singapore citizen; eligible for 12-month J-1 visa sponsored by NUS; program hosted by Stanford University.

AI Engineer & Founder specializing in high-performance agentic systems (MCP). Reduced delivery timelines from weeks to days (Tweakler) and documentation overhead by 78% (AssessMate).

Education

National University of Singapore

2025 – 2029

B.Comp. (Hons) in Computer Science • **Teaching Assistant, CS1101S Programming Methodology**

AIE Scholarship – Selected as 1 of 20 recipients

Technical Skills

Languages/Runtime: Python, Zig, C++, Rust, TypeScript/JavaScript, Node.js

AI & Data: MCP (stdio/HTTP), PyTorch, Transformers, OpenAI API, pgvector, FastAPI, Pydantic

Systems & Infrastructure: Linux, Docker, GitHub Actions CI/CD, Fly.io, PostgreSQL, pytest, Playwright

Experience

Founder, Tweakler

2026 – Present

- Founded Tweakler, a reference-led site generator: transforms one-line briefs into editable, brand-specific sites by deriving audience values and surfacing diverse design references.
- Reduced website delivery timelines from 3 - 4 weeks to under 7 days for studios like Brightsea Studio by building a reference-led generation system with an MCP workflow.

Research Assistant, National University of Singapore

Jan 2026 – Present

- Systematized the lab's evaluation of 3DGS methods by building a dynamic benchmark with 11 YAML pipelines and a 6-metric evaluator, enabling consistent performance tracking across 6 method repositories.

CCSGP Fellow (AssessMate), NUS Centre for Computing for Social Good

2025 – Present

- Reduced case-note documentation time from 90 to 20 minutes per child (a 78% reduction) by building AssessMate, an AI copilot that processes PDPA-sensitive data for early childhood intervention assessors.

C++ and Python Specialist (Freelance), Outlier.ai – Remote

Dec 2024 – May 2025

- Evaluated C++/Python model outputs for correctness, edge-case coverage, and reasoning quality, contributing feedback for code-generation model improvement.

Projects

EvoServe – Harness Evaluation & Quality Control

2026 – Present

- Built EvoServe, a harness-quality and evaluation control plane for coding-agent systems: compares policy candidates, ingests captured runs, independently verifies evidence, and produces auditable reports for quality and regression review.
- Implemented a 10-command validate/replay/score/verify/report CLI with 36 JSON schemas and 24 Zig tests; fails closed on malformed, missing, ambiguous, mutated, or untrusted evidence while preserving explicit review before policy activation.

GPT-2 (124M) Recreation from Scratch

2024

- Recreated GPT-2 124M from scratch (tokenizer, transformer blocks, optimizer, mixed precision, and FineWeb Edu evaluation); trained on 10B tokens.

WikiBase – Source-Grounded Knowledge Platform

2026

- Built WikiBase, a student knowledge platform that turns PDFs and notes into cited Markdown wikis, grounded Q&A, and flashcards; architected its FastAPI/PostgreSQL/pgvector backend, idempotent 5-stage pipeline, Docker/Fly.io release infrastructure, and GitHub Actions CI/CD plus a fail-closed 487-test harness covering migrations, security scans, and smoke checks.

Competitions

Replit Agent Hackathon

2026

- 1st place: Built Remember, an AR memory assistant with 4 AI workflows for face recognition, speech transcription, LLM memory extraction, and voice Q&A over 128-float face embeddings to a persistent personal knowledge graph for context-aware recall across repeated encounters.

Google DeepMind × Cactus (YC S25) Global Hackathon

2026

- 2nd place: Built SecureClaw, a privacy router for 12 PII types with 3 routing modes, on-device/cloud inference, and a 30-case benchmark.

Referral for Pranav Pappu

Internship | January - December 2027

To the Hiring Team,

I am Rohan Bahl, Founding Engineer at Cortex AI (YC F25) and I am pleased to strongly recommend Pranav Pappu for a January - December 2027 internship through the NUS Overseas Colleges Silicon Valley program. I have known Pranav through NUS hackathons, the AIE Hackathon, and several AI Tinkerers meetups. My recommendation is based on repeated technical discussions at such events, my review of EvoServe and Tweakler, and our conversations about their architecture and design decisions. These interactions have given me a clear view of how he frames engineering problems, explains trade-offs, and responds when his assumptions are examined closely.

EvoServe is Pranav's harness-quality/evaluation control plane for coding-agent systems. It is designed to compare policy candidates, review agent-generated artifacts, and improve harness performance over time through structured evaluation and evolutionary search. The design choice that stood out to me is that the producing harness is not allowed to certify its own result: EvoServe treats its success signals as untrusted, independently reruns deterministic checks, and verifies the score and execution context before reporting a result. In reviewing the system and discussing its architecture with him, I was impressed by the maturity of his design judgement. Pranav was not satisfied with producing an evaluation score; he focused on whether that score deserved to be trusted, what evidence supported it, and where uncertainty or failure should be surfaced.

Tweakler is his agent-driven system for using a selected permissioned website reference as a design foundation, then transforming it into a distinct, brand-oriented site. Its deterministic clone engine captures the reference with Playwright and reconstructs it as editable React code; agents then adapt the cloned structure, copy, and visual direction for a new audience while replacing the source imagery and other assets with approved or generated alternatives. Rather than asking an agent to invent a website from scratch, Tweakler gives it a strong visual foundation and validates the adapted result across desktop and mobile. Together, EvoServe and Tweakler show that Pranav uses AI for ambitious automation while placing deterministic checks and explicit trust boundaries around the results.

What distinguishes Pranav is that he looks beyond the happy path. He is a fast builder who turns ambitious ideas into working systems quickly without sacrificing code quality. In the work I reviewed, he maintained clear system boundaries, explicit verification paths, and careful implementation details. He engaged seriously with criticism and was willing to revise his reasoning when a stronger approach emerged. He combines the independence to build substantial systems with the judgement to ask how they fail and how their outputs can be verified.

Based on the work I reviewed and our repeated interactions across Singapore's AI builder community, I strongly recommend Pranav for an internship. He would be particularly well suited to AI infrastructure, evaluation, agent systems, and developer-tooling teams that value independent execution and technical rigor. I believe he would contribute meaningfully while learning quickly from the engineers around him.

Please feel free to contact me if further information would be helpful.

Sincerely,



Signature

05/08/2026

Date

Rohan Bahl

Founding Engineer, Cortex AI YC F25